



BSV Training

Eg06: Mergesort

An IP block* that sorts a vector in memory, using the “mergesort” algorithm. 06a uses an *ad hoc* testbench; 06b generalizes this to an “SoC” structure; 06c adds more parallelism to the mergesort.

algorithm. 06a

“SoC” struc

```

Import PkFrac*:
typedef BitVec{N} = Int{N};

module of_int_uint_half{N}:

  Integer file_depth = 32;

  function BitVec()  distroyed_gauss(BitVecT);
  return {x{N}};
endfunction


PFRac{N}(BitVecT) internal{N}
  *Note: the term
  *non-trivial hard
  *as units. Nowa
  returned PFRac{N}(file_depth) the...
PFRac{N}(BitVecT) outbound{N}
  returned PFRac{N}(file_depth) the...
PFRac{N}(BitVecT) internal{N}
  returned PFRac{N}(file_depth) the...

rule eval {True}:
  BitVecT in_data = bitvec{file_depth};
  PFRac{N}(BitVecT) out_gauss =
    distroyed_gauss(in_data
    out_gauss{N}(in_data );
  return out_gauss;
endrule : eval
endmodule : of_int_uint_half

```

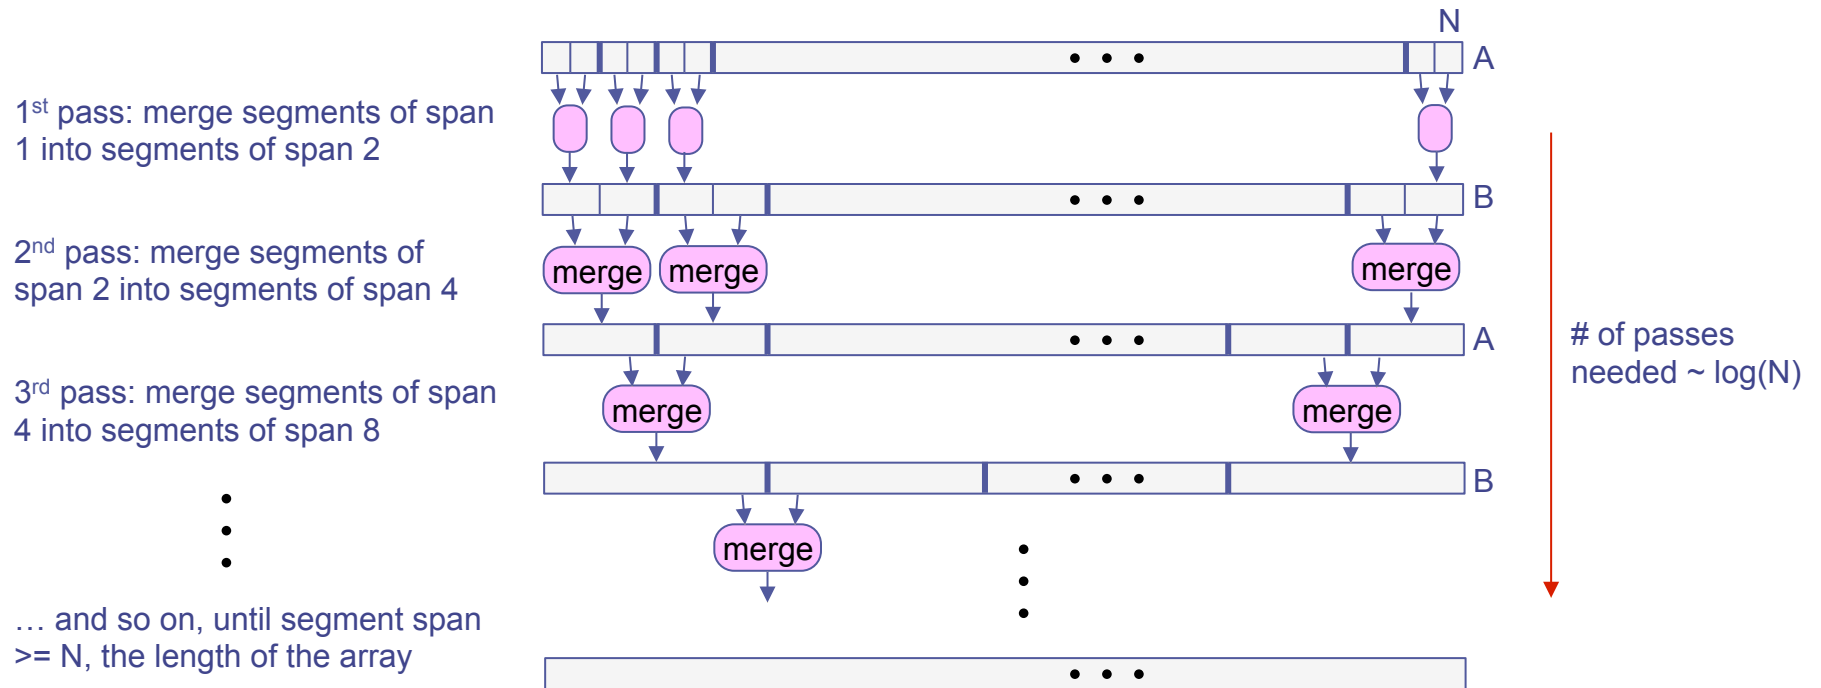
*Note: the term “IP block” or “Intellectual Property Block” originates from its reference to non-trivial hardware designs proprietary to a company and which are bought and sold as units. Nowadays the term is also used to refer to any non-trivial hardware block.

www.bluespec.com



Mergesort algorithm

Binary mergesort is a standard sorting algorithm, described in many textbooks and courses on algorithms. The basic idea is illustrated below.



Note: every segment that is input to “merge” is already sorted


Some edge conditions we need to take care of:

- N is usually not a power of 2, so last two spans may have unequal length
- Depending on N , the final sorted array may be in B , and so may have to be copied back into the original array A

If we complete one pass before starting the next, we can alternate between two arrays A and B .

Mergesort algorithm (contd.)

The “merge” step sorts two already-sorted segments of length ‘span’ into a sorted segment of length ‘2 x span’



```
merge (A,B,i,span) {  
    i0lim = i + span; i1lim = i + 2*span;  
    j = i0 = i; i1 = i0lim;  
  
    while (j < i1lim) {  
        if (i0 >= i0lim)          y = A[i1++]; // segment 0 exhausted; take from segment 1  
        else if (i1 >= i1lim)     y = A[i0++]; // segment 1 exhausted; take from segment 0  
        else if (A[i0] <= A[i1]) y = A[i0++]; // take from segment 0  
        else                     y = A[i1++]; // take from segment 1  
        B[j++] = y;  
    }  
}
```

Example variations


The accompanying code demonstrates three variations:

Eg06a_Mergesort/	Mergesort using a single “merge engine”. Testbench connects this to a single-port memory model.
Eg06b_Mergesort/	Generalizes system structure to an “SoC” containing an “CPU” (here, just an FSM), the mergesort IP module as a programmable “accelerator”, the memory model, and an interconnect fabric.
Eg06c_Mergesort/	Generalizes mergesort model to n parallel merge engines using n memory ports. Generalizes memory model to n memory ports. Handles memory ordering issue.

1st version: directory Eg06a_Mergesort/

Rather than create an *ad hoc* interface for our mergesort module, let us prepare it to be ready for plugging into an “SoC” (System on a Chip) as an “accelerator” module, illustrated to the right.

An SoC typically consists of CPUs, memories, an interconnect, and custom IP blocks (“Intellectual Property Blocks”) that perform particular functions for reasons of greater speed (acceleration) and/or less power consumption (compared to executing the same function in software on a CPU).



The interconnect fabric carries memory requests (red paths in the figure) and responses (green paths).

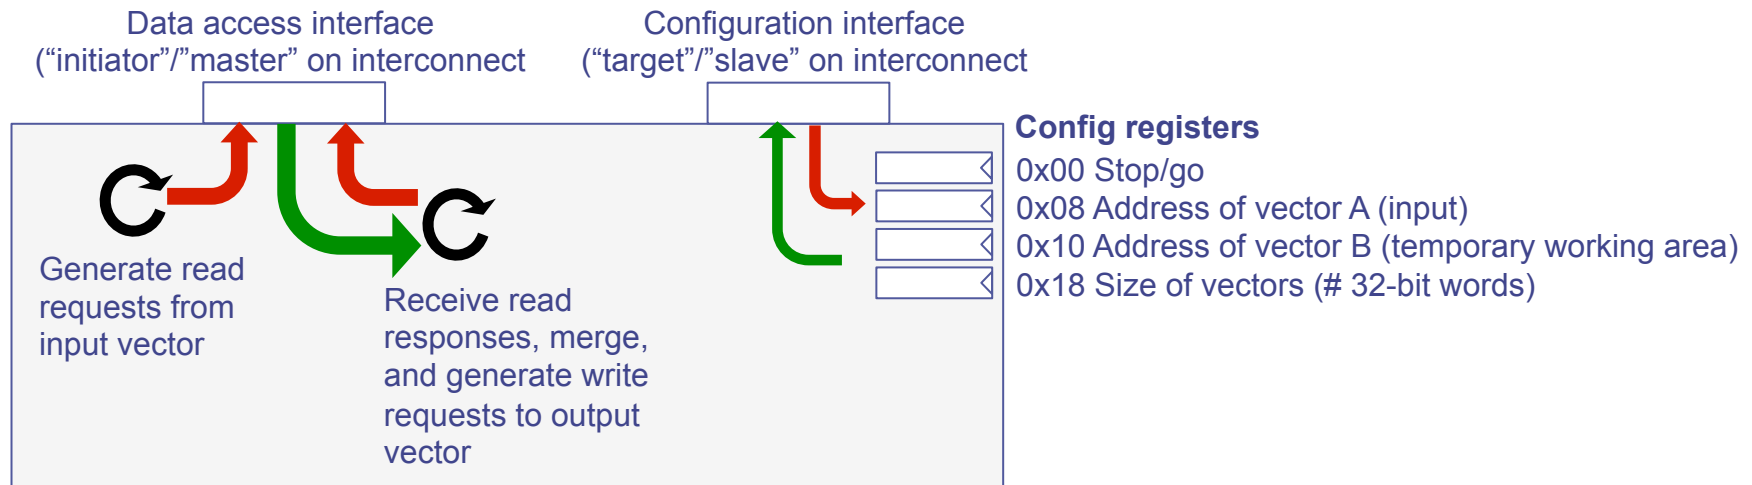
- *Initiator* ports (like the CPU port and the IP block read/write port) send requests and receive responses
- *Target* ports (like the Memory port and the IP block config port) receive requests and send responses

Memory requests are routed to the memory or to the IP block config port based on the address contained in the request. I.e., the (usually small number of) configuration registers in the IP block appear, to the CPU, just like memory locations at a particular base address (these addresses are disjoint from addresses serviced by the Memory block). We also say that the config registers are “memory-mapped”.

To operate the IP block, the CPU writes information needed for the operation to the config registers in the IP block, after which the IP block can perform its function (by reading and writing to memory). When the function is completed, the IP block may write a particular value to one of its config registers. The CPU can detect when the IP block has completed its function by “polling” (repeatedly reading) this register.

[In practice, IP blocks can also “interrupt” the CPU on completion; our examples here do not do this.]

1st version: directory Eg06a_Mergesort/ (contd.)



To perform the mergesort:

- The external environment must write the addresses and size of the vectors A and B to the config registers at offset 0x08, 0x10 and 0x18, and finally write a "1" (meaning: "start running") to the config register at offset 0
- The mergesort module then does its work, reading and writing through its data access port; when completed, it writes "0" to the config reg at offset 0
- The external environment can "poll" (repeated read) the config reg at offset 0, to detect completion

Time-out to reinforce some concepts

Please study the lectures:

- `Lec_Types` to review types, which are used to define memory requests and responses
- `Lec_Interfaces_TLM` to review the concepts behind interfaces like Get, Put, Client and Server, which are used for most of the interfaces in this example.
- `Lec_Interfaces_TLM` and `Lec_Typeclasses` to review the concepts behind the `mkConnection` abstraction, which is used in the testbench to connect all components together.
- `Lec_StmtFSM` for the concepts behind structured rule-based processes, which are used both in the mergesort module and in the testbench.
- `Lec_Interop_C` for the concepts behind importing C code into BSV, which is used in the memory model in this example.

Memory requests and responses: Common/Req_Rsp.bsv

```
typedef enum { READ, WRITE, UNKNOWN } TLMCommand
    deriving (Bits, Eq);

typedef enum { BITS8, ... BITS32, BITS64, ...} TLMBSize
    deriving (Bits, Eq);

typedef struct {
    TLMCommand      command;
    Bit #(addr_sz)  addr;
    Bit #(data_sz)  data;    // Only for write requests
    TLMBSize        b_size;
    Bit #(tid_sz)   tid;
} Req #(type tid_sz, type addr_sz, type data_sz)
    deriving (Bits);
```

Memory requests contain a command (READ/ WRITE), an address, data (for WRITE commands), a spec of the size of data being transferred, and a “transaction id” (tid).

```
typedef enum {OKAY, ..., SLVERR, DECERR} TLMStatus
    deriving (Eq, Bits);

typedef struct {
    TLMCommand      command;
    Bit #(data_sz)  data;
    TLMStatus        status;
    Bit #(tid_sz)   tid;
} Rsp #(type tid_sz, type data_sz)
    deriving (Bits);
```

Memory responses contain the original command (READ/WRITE), data (for READ commands), a status, and the original “transaction id” (tid).

Transaction Ids (tids) are common on memory requests/responses in modern SoCs, because:

- There may be multiple initiators, and the tid can serve as a “return address” identifying where a response should go
- Responses may be in a different order from the original requests, and the tid can identify the original order

1st version: directory Eg06a_Mergesort/ (contd.)

Specific choices for memory requests and responses in our Mergesort example

Common/Req_Rsp.bsv contains generic definitions for the types of memory requests and responses.

In particular, they are parameterized by the bit-widths of addresses (addr_sz), data (data_sz) and transaction ids (tid_sz), so that they can be used in various SoCs with various requirements.

In Eg06a_Mergesort/src_BSV/Sys_Configs.bsv, we make particular choices for these parameter for our mergesort example.

```
typedef 64 ASZ;
typedef 64 DSZ;

typedef Bit #(ASZ) Addr;
typedef Bit #(DSZ) Data;

typedef 1 TID_SZ_I;
typedef 1 TID_SZ_T;

typedef Bit #(TID_SZ_I) TID_I;
typedef Bit #(TID_SZ_T) TID_T;

typedef Req #(TID_SZ_I, ASZ, DSZ) Req_I;
typedef Rsp #(TID_SZ_I, DSZ) Rsp_I;

typedef Req #(TID_SZ_T, ASZ, DSZ) Req_T;
typedef Rsp #(TID_SZ_T, DSZ) Rsp_T;
```

Addresses are 64-bits wide

Data are 64-bits wide

Tids are 1-bit wide, both at initiators and targets

Req_I, Rsp_I, Req_T and Rsp_T are shorthand synonyms for requests and responses at initiators and targets with the specified sizes

1st version: directory Eg06a_Mergesort/ (contd.)

Specific choices for memory-mapping in our Mergesort example

In Eg06a_Mergesort/src_BSV/Sys_Configs.bsv, we also make particular choices for the number of memory ports and the addresses serviced by memory and the config registers

```
// Memory
Addr mem_base_addr = 0;
Addr mem_size      = 'h10_0000;
Addr mem_max_addr  = mem_base_addr + mem_size - 1;

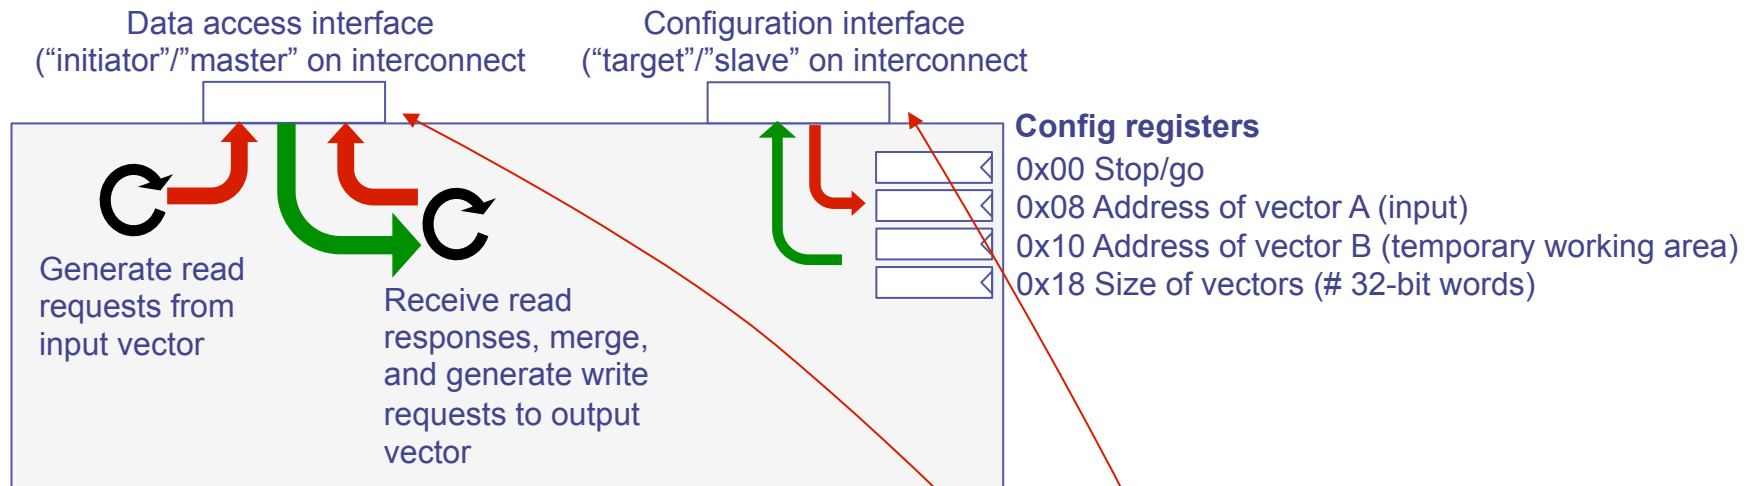
// Accelerator
Addr accel_base_addr = 'h80_0000;
typedef 4 N_Accel_Config_Regs;    // # of config registers
Addr accel_size      = fromInteger (valueOf (N_Accel_Config_Regs) * 8);
Addr accel_max_addr  = accel_base_addr + accel_size - 1;
```

Memory services addresses 0..0x10_0000-1

The accelerator (mergesort) has 4 config registers, each 8 bytes wide (64b), at base address 0x80_0000

1st version: directory Eg06a_Mergesort/ (contd.)

Interface for our mergesort module



In file Mergesort.bsv:

```
interface Mergesort_IFC;  
  method Action reset (Addr base_addr);  
  interface Server #(Req_T, Rsp_T) config_bus_ifc;  
  interface Client #(Req_I, Rsp_I) mem_bus_ifc;  
endinterface
```

1st version: directory Eg06a_Mergesort/ (contd.)

In Mergesort.bsv: mkMergeSort module structure

```
module mkMergeSort (Mergesort_IFC);  
  
  // -----  
  // Section: Configuration  
  ...  
  Vector #(N_Config_Regs, Reg #(Data)) vrg_configs;  
  ...  
  rule rl_handle_configReq;  
    ...  
  endrule  
  
  // -----  
  // Section: Merge sort behavior  
  ...  
  MergeEngine_IFC mergeEngine <- mkMergeEngine;  
  ...  
  mkAutoFSM (  
    seq  
      ...  
    endseq);  
  
  // -----  
  // INTERFACE  
  ...  
  interface mem_bus_ifc = mergeEngine.mem_bus_ifc;  
endmodule
```

Instantiate the configuration registers

This rule receives incoming config requests, reads/writes config regs, sends responses

Instantiate module for the “merge” step

This FSM implements the following pseudo-code:


```
while (True)  
  wait for 'run' command, init span=1, p1=A, p2=B  
  while (span < n) // do another pass:  
    i=0;  
    while (i < n)  
      merge (i, span, p1, p2);  
      i += 2*span;  
    swap p1,p2; span = 2x span  
  if final array is B, copy it back to A  
  config reg [0] = 0 (announce completion)
```

The mergeEngine's memory interface is directly used as the memory interface

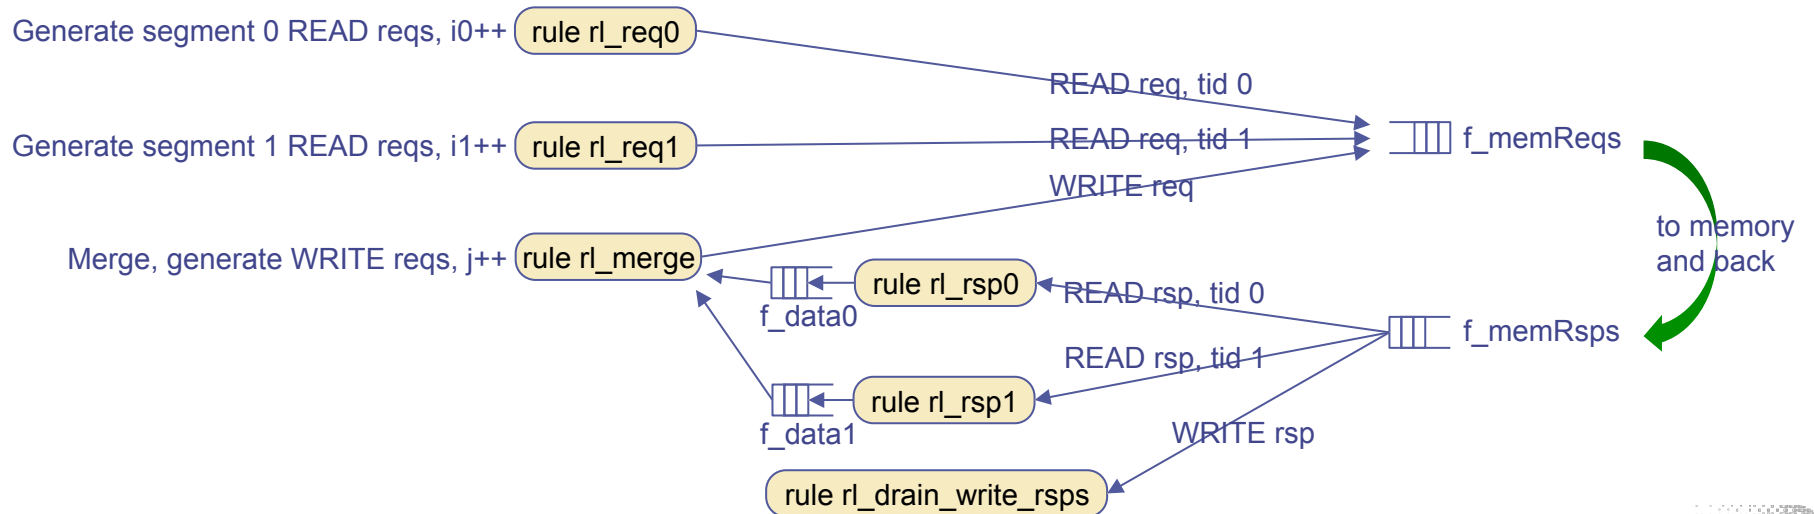
1st version: directory Eg06a_Mergesort/ (contd.)

In Mergesort.bsv: mkMergeEngine

This module implements the “merge” step which we saw earlier:



mkMergeEngine module data flow



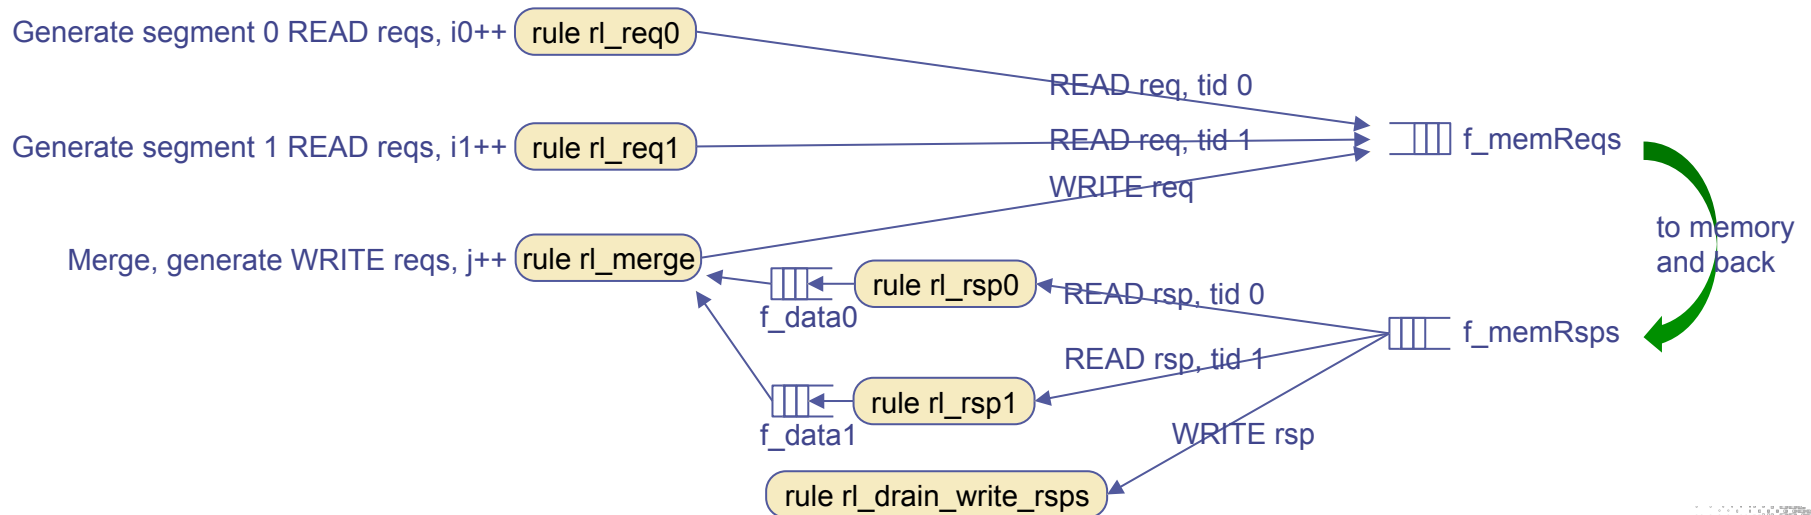
1st version: directory Eg06a_Mergesort/ (contd.)

mkMergeEngine is highly concurrent:

- rl_req0, rl_req1 and rl_merge continuously stream requests to memory
- rl_rsp0, rl_rsp1 and rl_drain... continuously handle the stream of responses

This is typical of high-performance accelerators which try to maximize utilization of available memory bandwidth. A software implementation on a CPU may not be able to generate such concurrent, pipelined memory accesses.

mkMergeEngine module data flow



1st version: directory Eg06a_Mergesort/ (contd.)

There is a danger of deadlock. Example:

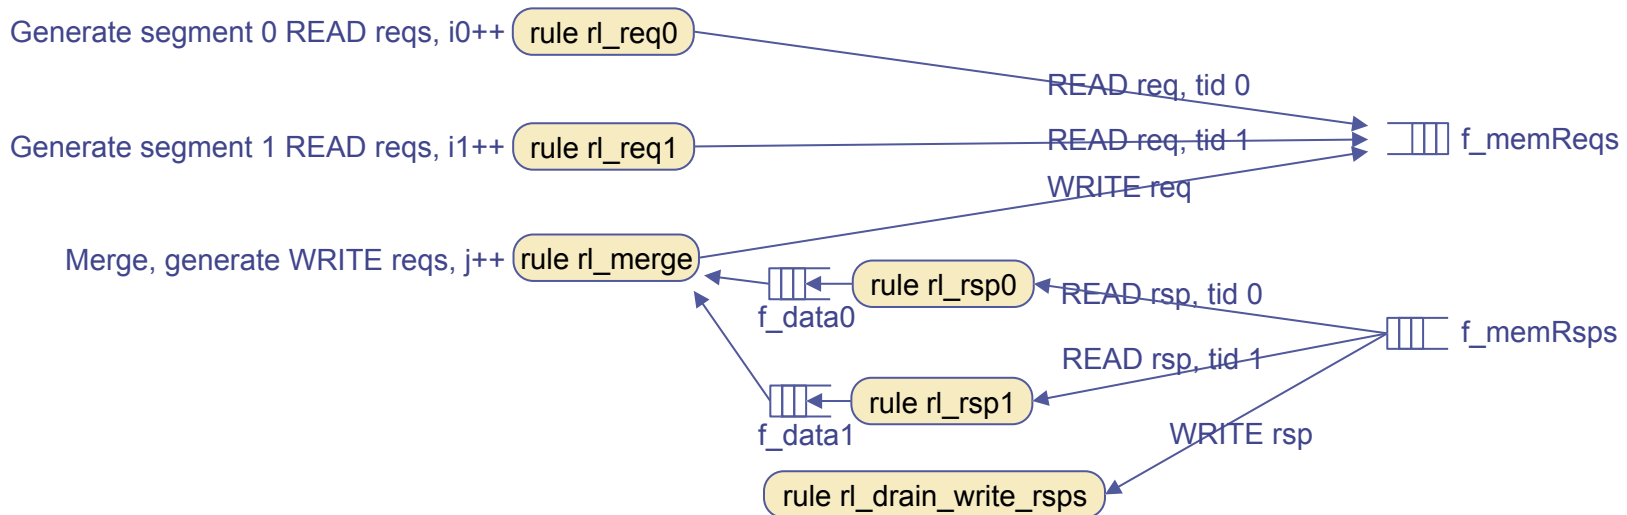
- Suppose `rl_merge` does not consume `f_data1`, for example because the next segment 1 item is > many segment 0 items
- Then, `f_data1` may become full, and if the first item in `f_memRsp`s is from segment 1, then we get stuck (the segment 0 items we need may be behind it). This kind of deadlock is called “head-of-line blocking”

Solution:

- The code has a parameter: `max_n_reqs_in_flight = 8`
- `f_data0` and `f_data1` are sized to accommodate 8 responses.
- `rl_req0` and `rl_rsp0` decrement and increment `ehr_credits0`, respectively, to never allow more than 8 requests in flight. `rl_req1` and `rl_rsp1` similarly maintain `ehr_credits1`.

This prevents the above deadlock situation.

mkMergeEngine module data flow



1st version: directory Eg06a_Mergesort/ (contd.)

In Common/Memory_Model.bsv: a memory model

To test our mergesort block, we need to provide it a memory containing the vector A to be sorted and the vector B for its scratch working area.

Large memories (particularly those implemented in DRAM) are typically not expressed in a hardware design language. Hence we merely use a *model* of memory for testing our IP block in simulation.

This is provided in Common/Memory_Model.bsv, which is excerpted below:

```
interface Memory_IFC;
  interface Vector #(N_Mem_Ports, Server #(Req_T, Rsp_T)) bus_ifc;
  ...
endinterface

module mkMemory_Model (Memory_IFC);
  ...
endmodule
```

The interface provides N_Mem_Ports servers for memory requests. In Eg06a we will only use 1 port, but in Eg06c we will increase this, to model a memory with higher bandwidth.

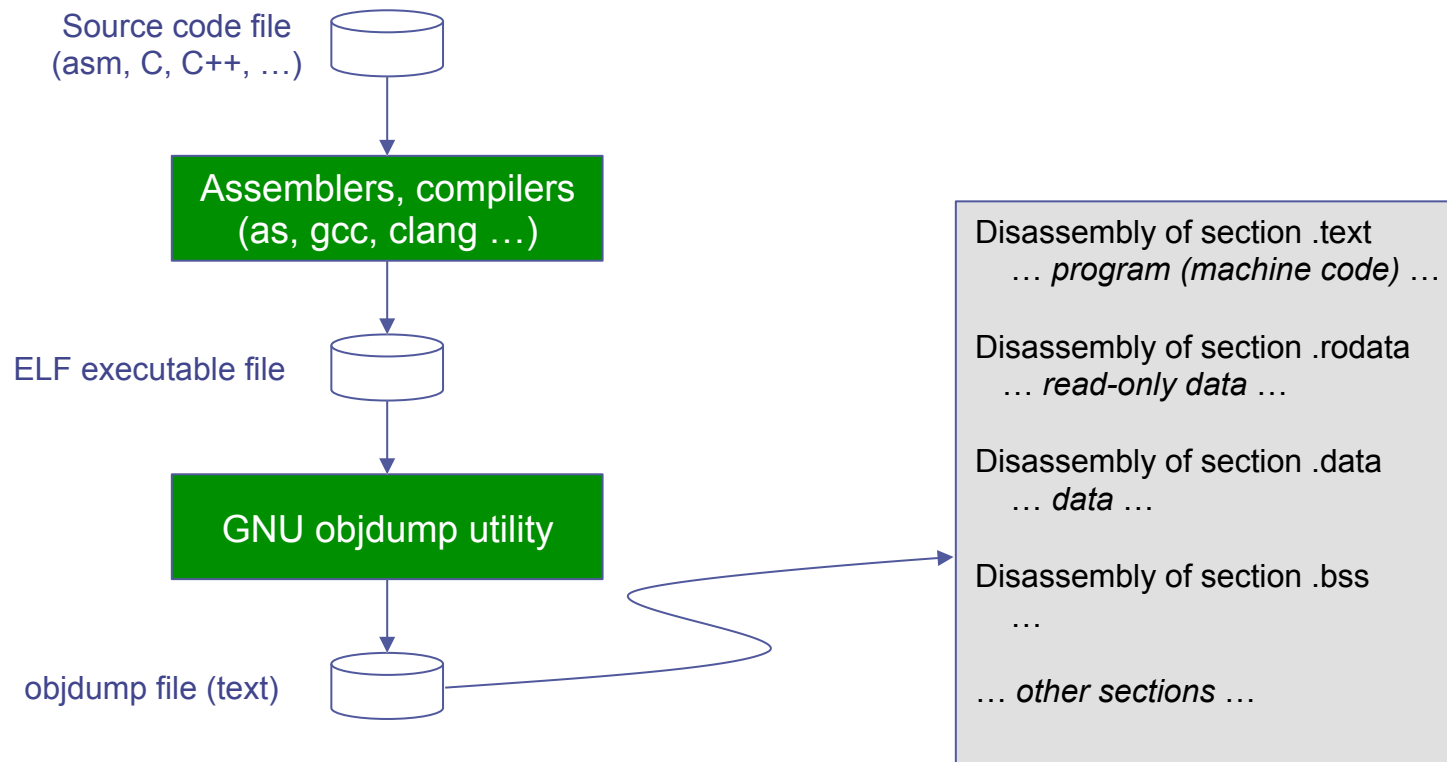
The body of the module mkMemory_Model is fairly straightforward. It illustrates how we can import a C function to do some work (in simulation only). The associated files are also in Common/:

- C_imports.{h,c} C functions to 'malloc' an array representing memory, and to read/write the array
- C_import_decls.bsv BSV declarations to connect BSV to the C functions

A word about “objdump” files

Our memory model initializes memory by reading in a file called “objdump”.

Objdump files are a standard format on several flavors of Unix (including Linux and OS X)



Note: We have not included any tools to create objdump files. The standard way is to use GNU tools. The pre-built objdump files distributed with this training were created using some ad hoc tools at Bluespec.

1st version: directory Eg06a_Mergesort/ (contd.)

In Testbench.bsv: a testbench

Our testbench is excerpted below:

```
module mkTestbench (Empty) ;

    Memory_IFC      mem                <- mkMemory_Model;
    Mergesort_IFC    mergesort          <- mkMergesort;

    mkConnection (mergesort.mem_bus_ifc, mem.bus_ifc [0]);
    ...
    mkAutoFSM (
        seq
            mem.initialize (mem_base_addr, mem_size, init_from_file);
            mergesort.reset (accel_base_addr);
            dump_mem_range;
            ... write mergesort's config regs ...
            ... loop, polling mergesort's config reg for completion ...
            dump_mem_range;
        endseq) ;
    endmodule
```

The testbench only performs a very small test so that you can easily inspect the outputs:

- Addr of the data array: 0x1000; scratch array: 0x1800; number of elements: 13

The 'dump_mem_range' calls (above) show memory contents before and after the sort.

(You are free to edit the program to try larger examples.)

Build and run the 1st version


- In the Build directory, build and run using the 'make' commands, with Bluesim and/or with Verilog sim, as described earlier
- Observe the inputs and outputs and verify that they are reasonable (final memory contents are a sorted version of initial memory contents)

2nd version: directory Eg06b_Mergesort/

In this version we use exactly the same Mergesort.bsv as in Eg06a.

We only generalize the environment around it into a “SoC” model.


Please study: `src_BSV/Sys_Configs.bsv`
in this directory, to see the changes to describe this new SoC.



Note that we do some “type-level” arithmetic to derive the number of fabric targets based on the number of memory ports; to derive “initiator numbers” (INums) based on the number of initiators, etc.:

```
typedef 1 N_Mem_Ports;  
typedef 2 Max_Initiators; // CPU Data access, Accelerator data port  
  
typedef TAdd #(N_Mem_Ports, 1) Max_Targets;  
  
typedef TLog #(Max_Initiators) INum_Sz;  
  
typedef Bit #(TLog #(Max_Initiators)) INum;  
typedef Bit #(TLog #(Max_Targets)) TNum;
```

2nd version: directory Eg06b_Mergesort/




Please study: `src_BSV/Sys_Configs.bsv`
in this directory, to see the changes to describe this new SoC.

Since we have more than one initiator on the fabric, transaction ids for targets have $\log(N)$ more bits than transaction ids for initiators, where N is the number of initiators, because the fabric must tack on $\log(N)$ bits to remember which initiator must get the corresponding response:

```
typedef 1 TID_SZ_I;  
  
// Transaction ids at targets  
typedef TAdd #(TLog #(Max_Initiators), TID_SZ_I) TID_SZ_T;  
  
typedef Bit #(TID_SZ_I) TID_I;  
typedef Bit #(TID_SZ_T) TID_T;
```

Finally, the end of the file now contains an “address decoder” module, which will be used by the Fabric to route memory requests either to the Memory or to the IP block, depending on the address.

2nd version: directory Eg06b_Mergesort/




Please study: `src_BSV/CPU.bsv`

You will see that it is not really a CPU, but just the testbench FSM from Eg06a, wrapped in a module that pretends to be a CPU. The interface is a step towards a real CPU interface, containing a “data-cache” interface, and other methods that will enable the GDB debugger to control the CPU:

```
interface CPU_IFC;
  // Interface to Data Memory
  interface Client #(Req_I, Rsp_I)  dcache_ifc;

  // GDB handling
  method Action      run_continue ();
  method CPU_Stop_Reason stop_reason;
  method Action      req_read_memW (Addr addr);
  method ActionValue #(Data) rsp_read_memW ();
  method Action      write_memW (Addr addr, Data d);
endinterface
```

2nd version: directory Eg06b_Mergesort/



Please study: [Common/Fabric.bsv](#)

It's interface is just a vector of Servers facing the initiators, and Clients facing the targets:

```
interface Fabric_IFC;
  interface Vector #(Max_Initiators, Server #(Req_I, Rsp_I)) v_servers;
  interface Vector #(Max_Targets, Client #(Req_T, Rsp_T)) v_clients;
endinterface
```

Recall that Req_I is different from Req_T, and Rsp_I is different from Rsp_T:

- For requests, the fabric will tack on extra transaction id (tid) bits to remember the “return address” where responses should go
- For responses, the fabric will use those extra tid bits to route the responses, and will also strip them off to restore the original tid.

The module mkFabric is quite straightforward. It represents a “full crossbar” switch, i.e., there is a separate datapath from each input port to each output port, in both directions. These are represented by the rules that are generated in for-loops. The only interesting thing in each rule is the tid manipulation as described above.

2nd version: directory Eg06b_Mergesort/


Please study: src_BSV/Testbench.bsv

The system is instantiated with this excerpt:

```
module mkTestbench (Empty) ;

  CPU_IFC      cpu      <- mkCPU_Model;
  Memory_IFC   mem      <- mkMemory_Model;
  Fabric_IFC   fabric   <- mkFabric;
  Mergesort_IFC mergesort <- mkMergesort;

  mkConnection (cpu.dcache_ifc, fabric.v_servers [cpu_d_iNum]);
  mkConnection (mergesort.mem_bus_ifc, fabric.v_servers [accel_iNum]);
  mkConnection (fabric.v_clients [mem_tNum], mem.bus_ifc [0]);
  mkConnection (fabric.v_clients [accel_tNum], mergesort.config_bus_ifc);
```



The behavior of the testbench (using mkAutoFSM) is a step towards a GDB-like interactive console to control the CPU. It uses an imported C command “c_console_get_command” to prompt the user for a GDB-like command and to return the command entered by the user, and then it executes the command.

In this first version, it recognizes just three commands: “continue” (to execute the program), “quit”, and “memory dump” to show a region of memory.

Build and run the 2nd version

- In the Build directory, build and run using the 'make' commands, with Bluesim and/or with Verilog sim, as described earlier
- When you simulate
 - You will initially see some INFO messages from the memory model, loading the objdump file
 - You will see the output of the first 'dumpmem', showing memory contents before the sort
 - Then, you will get a GDB-like prompt:

```
Command? [type 'h' for help]:
```

- Go ahead and type 'h' for a list of commands. It will list several GDB-like commands, but in this first version, it recognizes just three commands: "continue" (to execute the program), "quit", and "memory dump" to show a region of memory.
- Type 'c' for 'continue', and it will perform the mergesort and do the final 'dumpmem' showing the memory contents after sorting. You can also use the 'm' command to display this memory region.
- Type 'q' to quit back to the terminal prompt.

3rd version: directory Eg06c_Mergesort/

In this version we use the same source codes for the SoC environment (although we will increase the number of initiators and targets on the Fabric).

We generalize the Mergesort module to have multiple merge engines (instead of just one) that can operate in parallel. Each merge engine will have its own initiator port on the interconnect fabric. The source code is parameterized to have `N_Mergers`; in the example code we instantiate this to 2.

Correspondingly, we also increase the number of target ports for the memory. The source code is parameterized to have `N_Mem_Ports`; in the example code we set that to 2.


Please study: `src_BSV/Sys_Configs.bsv`

in this directory, to see the changes to describe this new SoC.

The overall structure is the same; just the details have changed to describe the extra target and initiator ports.

The address-decode function now interleaves memory addresses across the memory ports in 8-byte steps (addr 0..7 in first port, 8..15 in second port, ... and so on).


The file: `src_BSV/CPU.bsv`
is unchanged.



3rd version: directory Eg06c_Mergesort/ (contd.)

Please study: `src_BSV/Mergesort.bsv`

It's memory interface has now become a vector of Clients:



```
interface Mergesort_IFC;
  method Action reset (Addr base_addr);
  interface Server #(Req_T, Rsp_T) config_bus_ifc;
  interface Vector #(N_Mergers, Client #(Req_I, Rsp_I)) mem_bus_ifc;
endinterface
```

In the module `mkMergesort`, we now instantiate a vector of merge engines, instead of just one:

```
Vector #(N_Mergers, MergeEngine_IFC) mergeEngines <- replicateM (mkMergeEngine);
```

In the module `mkMergesort`'s behavior, where we used to start the single merge engine:

```
mergeEngine.start (0, 0, vrg_configs [n], rg_p1, rg_p2, vrg_configs [n]);
```


instead, we now just enqueue another "task" on to a queue:

```
f_tasks.enq (tuple4 (0, vrg_configs [n], rg_p1, rg_p2));
```

and, concurrently, we feed these tasks to any available merge engine:

```
for (Integer j = 0; j < valueOf (N_Mergers); j = j + 1)
  rule rl_exec_task;
    match { .i, .span, .p1, .p2 } = f_tasks.first; f_tasks.deq;
    mergeEngines[j].start (fromInteger (j), i, span, p1, p2, vrg_configs [n]);
  endrule
```

Memory ordering problem



When we introduce multiple memory ports with interleaved addresses in an SoC, we introduce a memory ordering problem:

- Suppose the IP block sends two requests, one after the other, to memory at two addresses Addr1 and Addr2.
- These requests may go to two different memory ports, depending on which port services which address
- Since those two ports may be access different regions of memory and face different delays and contention, the responses may come back in a different order
- If the IP block assumes that responses come back in the same order as requests, it will produce wrong results!

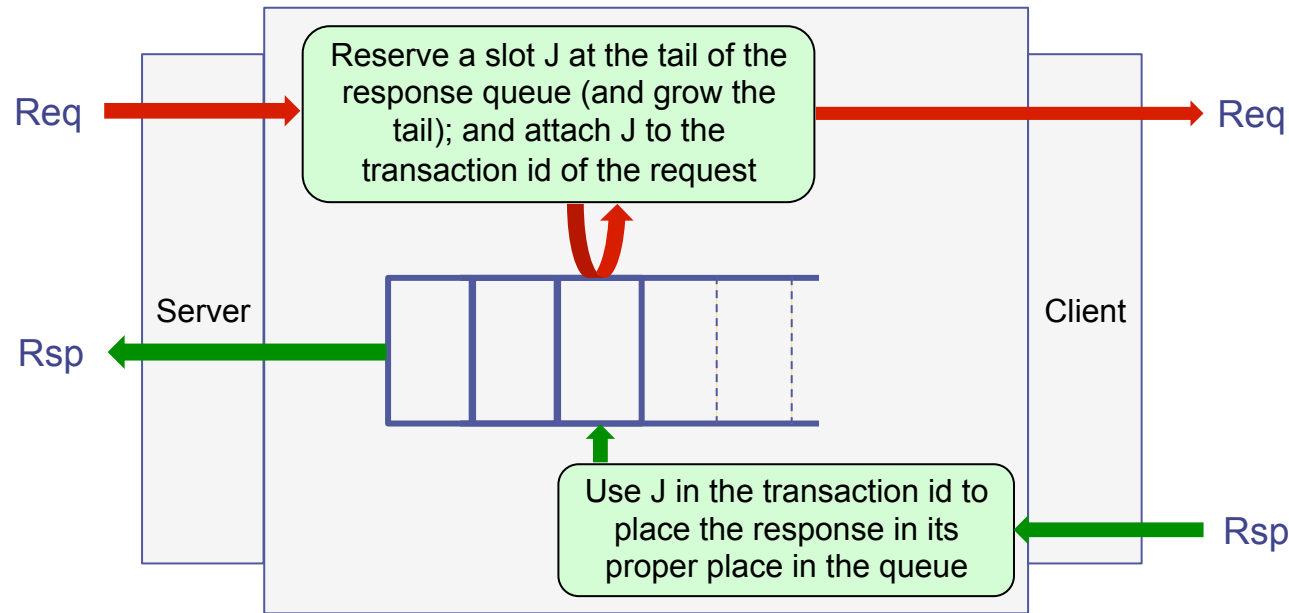

Our mkMergeSort module does assume that responses will come back in order!
(Exercise: study the code and convince yourself that it makes such an assumption.)

Solution: we place a “reorder buffer” between the IP block and the Fabric, to restore the order of responses delivered from the Fabric to the IP block.

3rd version: directory Eg06c_Mergesort/ (contd.)

Please study: Common/Reorder_Buffer.bsv

- An incoming request reserves a slot J at the current tail of the response queue (and grows the tail). The position J is carried along with the request in its transaction id.
- A response is inserted into its correct position in the response queue by looking at J in the transaction id



3rd version: directory Eg06c_Mergesort/ (contd.)

Please study: `src_BSV/Testbench.bsv`

The system is instantiated with this excerpt:

```
module mkTestbench (Empty) ;


  CPU_IFC      cpu      <- mkCPU_Model;
  Memory_IFC   mem      <- mkMemory_Model;
  Fabric_IFC   fabric   <- mkFabric;
  Mergesort_IFC mergesort <- mkMergesort;

  mkConnection (cpu.dcache_ifc, fabric.v_servers [cpu_d_iNum]);

  for (Integer j = 0; j < valueOf (N_Accel_Clients); j = j + 1) begin
    Reorder_Buffer_IFC reorder_buffer <- mkReorder_Buffer;
    mkConnection (mergesort.mem_bus_ifc [j], reorder_buffer.server);
    mkConnection (reorder_buffer.client, fabric.v_servers [accel_iNums[j]]);
  end

  for (Integer j = 0; j < valueOf (N_Mem_Ports); j = j + 1)
    mkConnection (fabric.v_clients [mem_tNums [j]], mem.bus_ifc [j]);

  mkConnection (fabric.v_clients [accel_tNum], mergesort.config_bus_ifc);
```



i.e., the for-loops connect the vector of mergesort initiators and memory targets to the corresponding ports of the interconnect fabric. We also place reorder buffers on the mergesort initiator ports.

Build and run the 3rd version

- In the Build directory, build and run using the 'make' commands, with Bluesim and/or with Verilog sim, as described earlier
- When you simulate you will see the same GDB-like command prompt as in Eg06b. Type 'c' (continue) to run the mergesort, and verify that the output looks reasonable.

Suggested exercises


- All three versions of the example sort 32-bit (4-byte) words of memory.
 - Modify the design to have a *static* parameter such that it will compile to a circuit that sorts memory in units of 1, 2, 4 or 8 bytes (static → the size is fixed at compile time). Note that Common/Req_Rsp.bsv already defines an enum type TLMBSize to specify byte size.
 - The mergesort engine should issue memory requests with the selected unit size.
 - Modify the design so that the byte-size selection is done *dynamically*:
 - Add another config register in which the CPU can specify the size.
 - The mergesort engine should issue memory requests with the selected unit size.
 - Modify the last design so that memory requests are always for 8 bytes, even if the sort is on smaller units. E.g., if the sort is on 1-byte units:
 - Only 1 memory read is needed to fetch 8 units.
 - Only 1 memory write is needed to store 8 units.
- All the examples perform a *binary* (radix 2) merge sort, i.e., the basic merge step merges *two* spans.
 - Modify the program to perform a radix 4 merge sort, i.e., the basic merge step should merge *four* spans. Question: when sorting an array of length n , how many memory references does this perform, compared to the binary merge sort?
 - Parameterize the module for a radix k mergesort, where k is a static parameter that may take some chosen range of values (2, 3, 4, ...).

Summary

This example has shown you key features of an IP block built for high-performance in an SoC context:

- Useful functionality (sorting, which is useful in *many* applications)
- Implementation using an efficient mathematical algorithm (mergesort)
- Key concepts of SoC structure: Fabrics, initiators, targets, memory mapping, ...
- Key concepts of high-performance: pipelining, task queue parallelism, memory bandwidth, managing out-of-order communication, ...
- Generality through parameterization on many dimensions (and hence capable of much re-use in other contexts)

bluespec



End

```

import P2P000;
typedef Bit[24] (outT);
module ex_hdl_csrR_hdl {empty};

Integer nfa_depth = 16;

function Bit[24] distribute_pump(outT out);
return (out);
endfunction

P2P000(outT) in_bounded;
in_bounded.P2P000(nfa_depth) the_in_bounded(in_bounded);
P2P000(outT) out_bounded;
out_bounded.P2P000(nfa_depth) the_out_bounded(out_bounded);
P2P000(outT) out_bounded;
out_bounded.P2P000(nfa_depth) the_out_bounded(out_bounded);

rule csp1 (True);
outT in_data = in_bounded.first;
P2P000(outT) out_data =
distribute_pump(in_data) == 0 ? out_bounded : out_bounded;
out_data.P2P000(out_data);
endrule;
endmodule : ex_hdl_csrR_hdl

```

