



# BSV Training

## Lec\_CRegs

A CReg (Concurrent Register) is a register-like primitive that enables greater concurrency (more rules per clock)

```

import PRCor*:
typedef Bit[24] (outT)

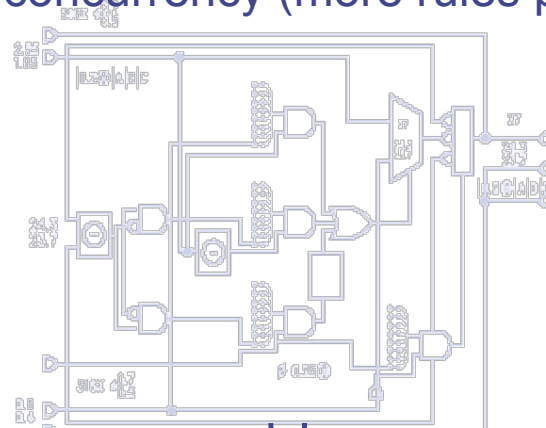
module ex_hdl_csrR_injEmpty()

  Integer nInDepth = 16


  function Bit[24] distribute_pump(outT out):
    return (out)
  endfunction

  PRCor(outT) inBoundC:
    outBoundC = PRCor(nInDepth) inBoundC
    PRCor(outT) outBoundC:
      outBoundC = PRCor(nInDepth) outBoundC
      PRCor(outT) outBoundC:
        outBoundC = PRCor(nInDepth) outBoundC
      PRCor(outT) outBoundC:
        outBoundC = PRCor(nInDepth) outBoundC
    endrule

  rule outC (True):
    outC = inBoundC
    PRCor(outT) outC:
      distribute_pump(outC) = 0 ? outBoundC : outBoundC
    endrule
endmodule
  
```




[www.bluespec.com](http://www.bluespec.com)



# Enabling greater concurrency

With a primitive module like `mkReg`, the effect of a rule's Action (`_write`) is not visible until the next clock (via `_read`), because of its method ordering constraint (`<`)



For greater rule concurrency, we need another primitive whose method ordering constraints allow an Action's effect to be visible in the same clock.

In BSV, we use a primitive called the *CReg*<sup>1</sup>.

<sup>1</sup> "CReg" = Concurrent Register. These are based on the "Ephemeral History Register" which was researched by Daniel Rosenband at MIT in 2004.

# A motivating example

- Suppose we want to build a two-port, saturating, up/down counter of 4-bit signed integer values, with the following interface:

```
interface UpDownSatCounter_Ifc;  
  method ActionValue #(Int #(4)) countA (Int #(4) delta);  
  method ActionValue #(Int #(4)) countB (Int #(4) delta);  
endinterface
```

- The “two ports” are the two identical methods countA and countB
- A module implementing this interface has internal state holding the current value of the counter (Int #(4) type, so range is -8 to +7)
- When either method is called,
  - The internal state is incremented by delta (range: -8 to +7), but saturates at +7 on overflow and at -8 on underflow
  - The old value of the counter is returned as the result of the method

*Note: because of finite precision and saturation, “count” operations are not commutative like in conventional arithmetic; so, the order of these operations matters here!*

# An implementation using ordinary registers (v1)

```
module mkUpDownSatCounter (UpDownSatCounter_Ifc);
  Reg #(Int #(4)) ctr <- mkReg (0);

  function ActionValue #(Int #(4)) fn_count (Int #(4) delta);
    actionvalue
      // Extend the precision to avoid over/under flows
      Int #(5) new_val = extend (ctr) + extend (delta);
      if (new_val > 7) ctr <= 7;
      else if (new_val < -8) ctr <= -8;
      else ctr <= truncate (new_val);

      return ctr;    // note: returns old value
    endactionvalue
  endfunction

  method countA (Int #(4) deltaA) = fn_count (deltaA);
  method countB (Int #(4) deltaB) = fn_count (deltaB);
endmodule
```

Since both methods do the same thing, we abstract their common behavior into a function `fn_count()`

## BSV notes:

- “extend (e)” sign-extends for `Int#(n)`, and zero-extends for `Bit#(n)` and `UInt#(n)`
- “truncate (e)” drops MSBs, taking care of sign bits etc.
- The number of bits extended/truncated depends on the input and output type widths

# A testbench to drive the up/down counter module

```
module mkTest (Empty);
  UpDownSatCounter_ifc ctr <- mkUpDownSatCounter;
  Reg #(int) step <- mkReg (0);
  Reg #(Bool) flag0 <- mkReg (False); Reg #(Bool) flag1 <- mkReg (False);

  function Action count_show (Integer rulenum, Bool a_not_b, Int #(4) delta);
    action
      let x <- (a_not_b ? ctr.countA (delta) : ctr.countB (delta));
      $display ("cycle %0d, r%0d: is %0d, count (%0d)", cur_cycle, rulenum, x, delta);
    endaction
  endfunction

  // Rules 0-9 are sequential, just testing one method at a time
  rule r0 (step == 0); count_show (0, True, 3); step <= 1; endrule
  rule r1 (step == 1); count_show (1, True, 3); step <= 2; endrule
  ... and similarly, sequentially feed deltas of 3,3, -6,-6,-6,-6, 7, 3,
  // Concurrent execution
  rule r10 (step == 10 && !flag0); count_show (10,True, 6); flag0 <= True; endrule
  rule r11 (step == 10 && !flag1); count_show (11,False, -3); flag1 <= True; endrule

  // Show final value
  rule r12 (step == 10 && flag0 && flag1); count_show (12,True, 0); $finish; endrule
endmodule: mkTest
```

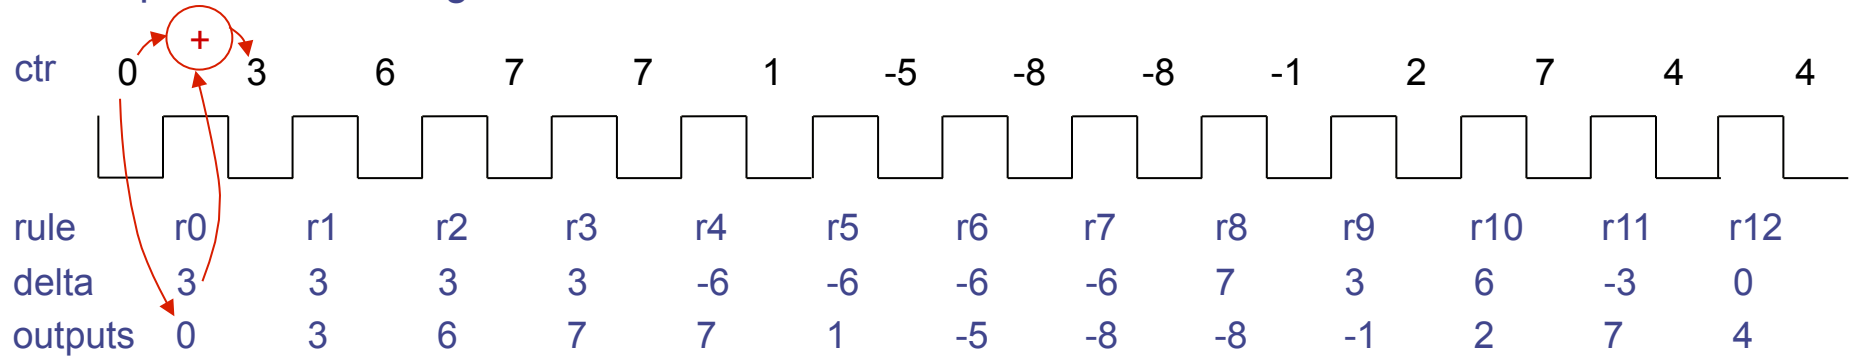
In rules 0-9, we call either countA or countB with deltas: 3,3,3,3, -6,-6,-6,-6, 7, 3  
The rule conditions and step assignments force them to fire 1 rule per clock (and so it doesn't matter whether we call countA or countB in these rules).

Rules 10 and 11 could potentially fire concurrently (if scheduling permits).

Rule 12 just displays the final counter value and exits.

# Expected behavior and outputs for v1

We expect the following behavior if r10 fires in the clock before r11:



We expect the following behavior if r11 fires in the clock before r10:



# Actual output for v1

When we compile the program (v1), *bsc* produces the following message:

```
Warning: "Test.bsv", line 16, column 8: (G0010)
Rule "r10" was treated as more urgent than "r11". Conflicts:
  "r10" cannot fire before "r11": calls to ctr.countA vs. ctr.countB
  "r11" cannot fire before "r10": calls to ctr.countB vs. ctr.countA
```

This is saying:

- r10 and r11 conflict; they cannot be scheduled in the same clock  
(countA and countB conflict because they both read and write the “ctr” register inside mkUpDownSatCounter, thus both possible rule orders will violate a “\_read < \_write” ordering constraint)
- *bsc* has chosen to give priority to r10, i.e., if both r10 and r11 are enabled in the same clock, the scheduling logic will allow r10 to fire and will suppress r11  
(r11 could fire, and indeed it does, in the next clock, when r10 is no longer enabled)
- Note: you can force the opposite priority by adding a “descending\_urgency” attribute to the module

When we run the program (v1), we see:  
(per first schedule in previous slide)

```
cycle 1, r0: is 0, count (3)
cycle 2, r1: is 3, count (3)
cycle 3, r2: is 6, count (3)
cycle 4, r3: is 7, count (3)
cycle 5, r4: is 7, count (-6)
cycle 6, r5: is 1, count (-6)
cycle 7, r6: is -5, count (-6)
cycle 8, r7: is -8, count (-6)
cycle 9, r8: is -8, count (7)
cycle 10, r9: is -1, count (3)
cycle 11, r10: is 2, count (6)
cycle 12, r11: is 7, count (-3)
cycle 13, r12: is 4, count (0)
```

2 cycles

## v1 is not really a “2-port” counter

v1 of our mkUpDownSatCounter may be functionally correct, but it’s hardly a “2-port” counter!

When we say “2-port”, we are making a performance characterization, i.e., we expect both ports to be operable in the same clock.

For this, we need to replace the Reg in mkUpDownSatCounter with an CReg, a different primitive that allows “multiple reads and writes” within a clock.



# First: specifying the semantics of the two ports

Before we worry about implementations and CRegs, we must first specify the *desired semantics* of the two ports! Specifically:

When both countA and countB are operated in the same clock,

- what should be the final value of the counter?
- what should be the “old” values returned by each method?

In light of the finite precision arithmetic, and the saturating behavior, there is no obvious unique answer! It is a design choice!

In RTL designs, this is typically where you’ll see an *ad hoc* choice made by the designer

- Which is (hopefully!) implemented correctly
- Which is (hopefully!) documented clearly and fully in English text in the datasheet
- Which may contain usage rules the user of the IP must follow, and which therefore need verification

In BSV, method orderings give us a formal and precise way to specify the semantics. By specifying that we want “countA < countB” or “countA > countB”, we give precise answers to the above two semantic questions, because when operated in the same clock, there is a well-defined *logical* ordering that specifies the behavior exactly.

Further, *bsc* always verifies correct usage because it’s in the semantics, not *ad hoc* English.

# CRegs (Concurrent Registers)

A CReg provides a *vector* of standard Reg interfaces that can be operated concurrently:

```
interface CReg #(numeric type n, type t);  
    interface Vector #(n, Reg #(t)) ports;  
endinterface
```

BSV notes:

- “Vector” is a standard importable BSV library
- The parameter n is the number of elements in the vector; t is the data type stored in the CReg

The ports of an CReg can be operated concurrently, with the following ordering constraints:

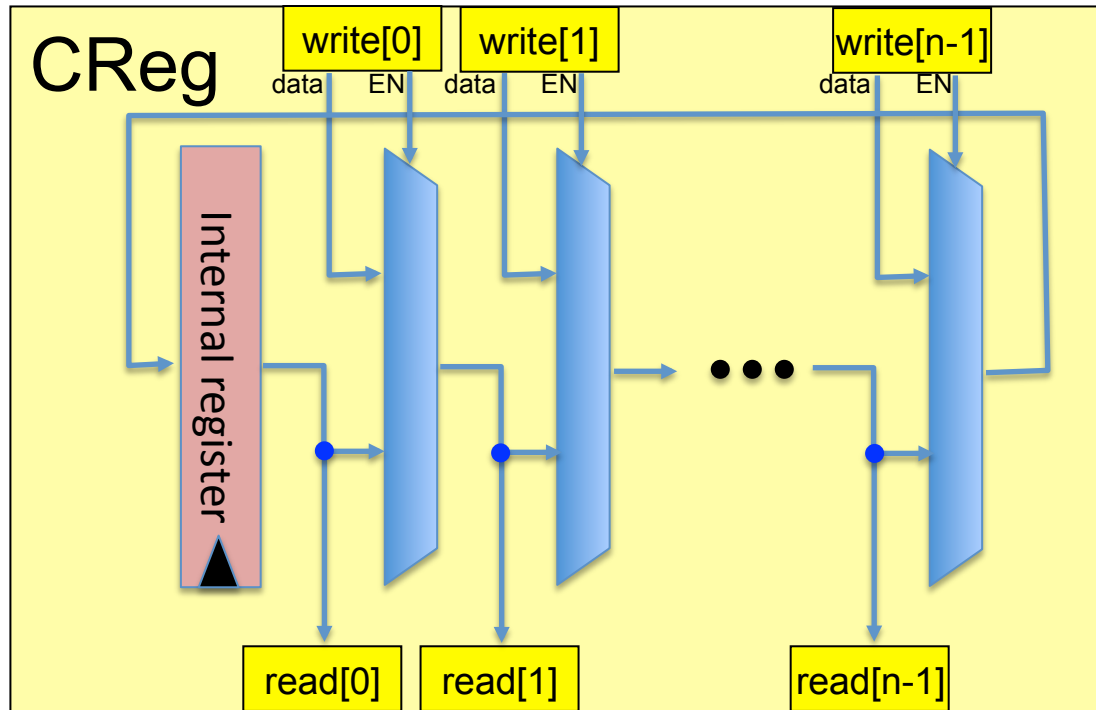
```
ports [0]._read <= ports [0]._write <  
ports[1]._read <= ports[1]._write <  
ports[2]._read <= ports[2]._write <  
... ..  
ports[n-1]._read <= ports[n-1]._write
```

*This is the same as the standard register method-ordering constraint*

*But note that a value written in port 0 can be read concurrently on port 1 (by a logically later rule in the same clock), unlike an ordinary register where a write can only be read in the next clock*

# A possible implementation of an CReg

This figure shows a possible circuit implementation of a CReg:



But note, this is not a *definition* of an CReg, it is merely shown to strengthen intuition. It is important, as usual, to keep separate the logical semantics from any implementation semantics. When using CRegs in BSV, one only needs to consider its method-ordering constraints (shown in the previous slide).

# Implementing our counter using CRegs (v2)

```
module mkUpDownSatCounter (UpDownSatCounter_Ifc);
  CReg#(2, Int #(4)) ctr <- mkCReg(0); // 2 ports

  function ActionValue #(Int #(4)) fn_count (Integer p, Int #(4) delta);
    actionvalue
      // Extend the precision to avoid over/under flows
      Int #(5) new_val = extend (ctr.ports [p]) + extend (delta);
      if (new_val > 7) ctr.ports [p] <= 7;
      else if (new_val < -8) ctr.ports [p] <= -8;
      else ctr.ports [p] <= truncate (new_val);

      return ctr.ports [p]; // note: returns old value
    endactionvalue
  endfunction

  method countA (Int #(4) delta) = fn_count (0, delta);
  method countB (Int #(4) delta) = fn_count (1, delta);
endmodule
```

Change Reg to  
CReg

Add CReg port  
selections to  
reads and writes

For “countA < countB”.

To implement “countB < countA”,  
change to:


```
... = fn_count (1, delta);
... = fn_count (0, delta);
```

This is only a slight change to v1:

- The internal “ctr” is now a 2-port CReg instead of a Reg
- fn\_count is now parameterized by the CReg port “p” it should use
- countA and countB call this function with ports 0 and 1, respectively, thereby implementing the ordering semantics “countA < countB”

# Behavior and outputs for v2

We expect the following behavior ( $r_{10} < r_{11}$  in same clock):




When we run the program (v2), we see:

```
cycle 1, r0: is 0, count (3)
cycle 2, r1: is 3, count (3)
cycle 3, r2: is 6, count (3)
cycle 4, r3: is 7, count (3)
cycle 5, r4: is 7, count (-6)
cycle 6, r5: is 1, count (-6)
cycle 7, r6: is -5, count (-6)
cycle 8, r7: is -8, count (-6)
cycle 9, r8: is -8, count (7)
cycle 10, r9: is -1, count (3)
cycle 11, r10: is 2, count (6)
cycle 11, r11: is 7, count (-3)
cycle 12, r12: is 4, count (0)
```

*same cycle*

## A second example

Consider a 2-stage CPU pipeline:



Let us focus on the FIFO connecting the stages:

Usually this is just a 1-element FIFO (we call it a “PipelineFIFO”).

a.k.a. “pipeline register with interlock” (the interlock is just the extra valid bit that allows the execute stage to stall if there is nothing in the pipeline register, and allows the fetch/decode stage to stall if there is already something in the register which has not been consumed by the execute stage).

# An implementation using ordinary registers (v1)

```
module mkFIFO1 (FIFO #(t));  
  Reg #(t)      rg      <- mkRegU;      // data storage  
  Reg #(Bit #(1)) rg_count <- mkReg (0); // # of items in FIFO (0 or 1)  
  
  method Bool notEmpty = (rg_count == 1);  
  method Bool notFull  = (rg_count == 0);  
  
  method Action enq (t x) if (rg_count == 0); // can enq if not full  
    rg <= x;  
    rg_count <= 1;  
  endmethod  
  
  method t first () if (rg_count == 1); // can see first if not empty  
    return rg;  
  endmethod  
  
  method Action deq () if (rg_count == 1); // can deq if not empty  
    rg_count <= 0;  
  endmethod  
  
  method Action clear;  
    rg_count <= 0;  
  endmethod  
endmodule
```

But: enq and {first, deq} could never be concurrent, with mutually exclusive conditions: `rg_count == 0` and `rg_count == 1`

Implication → the fetch/decode stage and the execute stage in the 2-stage CPU pipeline could never execute in the same clock (it isn't really a pipeline!)

# First: specify desired semantics of concurrent methods

Before we worry about implementations, we must first specify the desired *semantics* of concurrency on FIFO methods. In BSV we commonly use the following two kinds of FIFOs:

## PipelineFIFOs:

- When empty, only enq is enabled
- When full, enq, first and deq are enabled, with:  $\{\text{first}, \text{deq}\} < \text{enq}$   
i.e., if both methods are enabled, logically it is like  $\{\text{first}, \text{deq}\}$  followed by enq,  
i.e., data currently in the FIFO is returned for  $\{\text{first}, \text{deq}\}$ , and new data is enqueued.

## BypassFIFOs:

- When full, only  $\{\text{first}, \text{deq}\}$  is enabled
- When empty, enq, first and deq are enabled, with:  $\text{enq} < \{\text{first}, \text{deq}\}$   
i.e., if both methods are enabled, logically it is like enq followed by  $\{\text{first}, \text{deq}\}$ ,  
i.e., the newly enqueued value is “bypassed” through to  $\{\text{first}, \text{deq}\}$ .



# An implementation of Pipeline FIFOs using CRegs

```
module mkPipelineFIFO (FIFO #(t));  
  CReg #(3, t)      crg      <- mkCRegU;      // data storage  
  CReg #(3, Bit #(1)) crg_count <- mkCReg (0); // # of items in FIFO  
  
  method Bool notEmpty = (crg_count.ports[0] == 1);  
  method Bool notFull  = (crg_count.ports[1] == 0);  
  
  method Action enq (t x) if (crg_count.ports[1] == 0);  
    crg.ports[1] <= x;  
    crg_count.ports[1] <= 1;  
  endmethod  
  
  method t first () if (crg_count.ports[0] == 1);  
    return crg.ports[0];  
  endmethod  
  
  method Action deq () if (crg_count.ports[0] == 1);  
    crg_count.ports[0] <= 0;  
  endmethod  
  
  method Action clear;  
    crg_count.ports[2] <= 0;  
  endmethod  
endmodule
```

This is only a slight change to v1:

- notEmpty, first and deq use CReg port 0
- notFull and enq use CReg port 1
- clear uses CReg port 2

# An implementation of BypassFIFOs using CRegs

```
module mkBypassFIFO (FIFO #(t));
  CReg #(3, t)      crg      <- mkCRegU;      // data storage
  CReg #(3, Bit #(1)) crg_count <- mkCReg (0); // # of items in FIFO

  method Bool notEmpty = (crg_count.ports[1] == 1);
  method Bool notFull  = (crg_count.ports[0] == 0);

  method Action enq (t x) if (crg_count.ports[0] == 0);
    crg.ports[0] <= x;
    crg_count.ports[0] <= 1;
  endmethod

  method t first () if (crg_count.ports[1] == 1);
    return crg.ports[1];
  endmethod

  method Action deq () if (crg_count.ports[1] == 1);
    crg_count.ports[1] <= 0;
  endmethod

  method Action clear;
    crg_count.ports[2] <= 0;
  endmethod
endmodule
```

This is only a slight change to v1:

- notFull and enq use CReg port 0
- notEmpty, first and deq use CReg port 1
- clear uses CReg port 2

# CReg summary

The CReg is a highly concurrent primitive, i.e., it has multiple methods that can be invoked by multiple rules within a clock in a well-defined logical sequential order.

When using a CReg to communicate between rules that you want to be concurrent (i.e., able to fire in the same clock),

- first, be clear about what semantics you want, by thinking about what logical ordering of rules you want
- then, use CRegs to implement that ordering
  - (ascending CReg port indexes directly correspond to ordering)

Note: a design using CRegs will be functionally correct with any schedule, even one rule per clock. In the extreme schedule of one rule per clock, an CReg is exactly equivalent to an ordinary register (using mkReg).

In practice, we more often directly use concurrent library modules like PipelineFIFO and BypassFIFO. If necessary, we use CRegs to implement a concurrent module that is not available in the library.

# Hands-on

- BSV-by-Example book: Examples in Chapter 8



# End

```
Export PFCoinType;

typedef struct PFCoinType;

module ex_hl_coint_hlcoinType;

Integer ffcointType;

function HlCoinType()
return (new PFCoinType());
endfunction

PFCoinType() hlcoinType;
return HlCoinType(ffcointType) the_hlcoinType hlcoinType;
PFCoinType() hlcoinType;
return HlCoinType(ffcointType) the_hlcoinType hlcoinType;
PFCoinType() hlcoinType;
return HlCoinType(ffcointType) the_hlcoinType hlcoinType;

endmodule

endmodule

endmodule
```

## Questions?

Join online forums at [www.bluespec.com](http://www.bluespec.com), and ask your question,  
or send an e-mail to [support@bluespec.com](mailto:support@bluespec.com)

